

Concurrency: Sequential execution limits resource use and performance. Concurrency interleaves execution of threads or processes to improve utilization and responsiveness.

- **Race Conditions:** Occur when threads access shared resources simultaneously, causing unpredictable results due to interleaved operations.
- **Critical Sections:** Code segments accessing shared resources, requiring protection to prevent race conditions.

Synchronization Primitives: Tools to coordinate concurrent threads, ensuring orderly access to shared resources and preventing race conditions and deadlocks.

- **Locks:** Basic mutual exclusion allowing only one thread in a critical section.
 - Issue: High contention leads to inefficient waiting.
- **Spin Locks:** Threads "spin," repeatedly checking until the lock is free. Use for short critical sections with low contention, as they waste CPU during high contention. Issue: Spinning wastes CPU resources under high contention.
 - Disable Interrupts: Prevents context switches on uniprocessor systems. Fixes context switching but monopolizes the CPU.
 - Load/Store: Loads the lock's current value to check if it's free (0), then stores 1 to claim it. Issue: Not atomic, so race conditions can occur if multiple threads load and store simultaneously.
 - Test-and-Set: Atomically sets the lock's value to 1 and loads the previous value. If the previous value was 0 (unlocked), the thread acquires the lock; otherwise, it retries. Issue: Causes repeated retries (busy-waiting) when locked.
 - Compare-and-Swap (CAS): Atomically loads the lock's current value and compares it with 0. If 0, CAS stores 1 to claim the lock. If not, CAS does nothing, and the thread retries. Issue: Still involves busy-waiting, which wastes CPU.
- **Blocking Locks** (e.g., Mutexes): Threads block (sleep) if the lock is unavailable, reducing CPU waste. Suitable for long critical sections with high contention. Issue: Can cause deadlocks if threads hold resources while waiting for others; lacks conditional waiting.
- **Condition Variables:** Used alongside blocking locks (like mutexes) to allow threads to wait for specific conditions and signal others when conditions change.
 - `wait(mutex)`: Releases the mutex and waits, reacquiring it on wakeup.
 - `signal() / broadcast()`: Wakes up one or all waiting threads.
 - Issue: Spurious wakeups require rechecking the condition; doesn't manage access to multiple resources.
- **Bounded Buffer with Condition Variable:** Uses two condition variables, `notEmpty` and `notFull`, for producer and consumer threads. Producers wait on `notFull` when the buffer is full; consumers wait on `notEmpty` when it's empty.
- **Producer/Consumer with Two Condition Variables:** `notEmpty` signals consumers when items are available; `notFull` signals producers when there's space.
- **Semaphores:** Control access to multiple instances of a resource. A non-negative integer count represents available resources.
 - `acquire()`: Decrements count if positive; blocks if zero.
 - `release()`: Increments count; may wake waiting threads.
 - **Bounded Buffer with Semaphores:** Use a semaphore empty for free slots and full for filled slots, alongside a mutex to control access. Producers wait(empty) before adding; consumers wait(full) before removing.
- **Equivalence:** Synchronization primitives are functionally equivalent by adapting usage. A binary semaphore acts as a mutex by setting its count to 0 (acquire) and 1 (release) for mutual exclusion. Conversely, a mutex and condition variable mimic a counting semaphore by using a resource count and signalling waiting threads, enabling controlled access.

Deadlock: Occurs when threads are blocked, each waiting for resources held by others. 4 Conditions must be met to reach a deadlock, if any one condition is prevented, the deadlock is broken:

1. Mutual Exclusion: Only one thread can use a resource at a time.
 2. Hold and Wait: Threads hold resources while waiting for others.
 3. No Pre-emption: Resources cannot be forcibly taken away.
 4. Circular Wait: Circular chain of threads each waiting for resources held by the next.
- Avoidance:**
- Wait-Free Algorithms: Ensure every operation completes in a finite number of steps, removing the need to wait.
 - Atomic Lock Acquisition: Require threads to acquire all needed resources atomically.
 - Pre-emption: Allow resources to be forcibly taken away.
 - Resource Ordering: Impose an order on resource acquisition to prevent circular wait.
 - **Banker's Algorithm:** Checks if granting a resource can be done safely without causing a deadlock. Parameters include:
 - Safe State: A state where the system can allocate resources to each process in some order and still avoid deadlock.
 - Available Resources: The number of available resources of each type.
 - Maximum Demand: The maximum number of resources a process may need.
 - Allocated Resources: The number of resources currently allocated to a process.
 - Need: The remaining resources that a process may still request.
 - a. Check Need: Ensure the request doesn't exceed what the process still needs.
 - b. Check Availability: Confirm enough resources are available to fulfill the request.
 - c. Simulate Allocation: Temporarily grant the resources and verify if the system stays safe.
 - d. Decision: If safe, grant the request. If not, deny and make the process wait.

Given: 3 people (Alice, Bob, Charlie), 3 fruits (Apples, Bananas, Oranges), Available: 2 1 0
 Alice requests 1 1 0 (Need: 2 2 0, Allocated: 1 0 1) → Granted, no one starves
 Charlie requests 1 0 1 (Need: 1 1 1, Allocated: 0 1 0) → Denied, exceeds Need

I/O Devices: Users need to be able interact with OS via external devices. OS Needs standardised way to interact with diverse hardware having different speeds and protocols.

- **Canonical devices** solve hardware diversity through standardized registers (status: reports device state, command: sends instructions, data: transfers information). Driver translates between standard interface and hardware. Device interaction happens through:
 - **Polling:** CPU repeatedly checks device status. Issue: Wastes CPU cycles - good for fast devices or predictable timing.
 - **Interrupts:** Device signals CPU when ready. Issue: Polling inefficient. Solution: CPU free to do other work while waiting.
 - **Direct Memory Access (DMA):** Issue: CPU wastes cycles copying data, so DMA transfers data between memory and devices independently. CPU only sets up transfer then continues work; gets interrupted when done.
- **Patterns:** What does the process do while waiting on I/O?
 - **Blocking:** Process waits until I/O completes, simple but idle CPU.
 - **Non Blocking:** Process starts I/O, keeps running, checks status periodically. Improves CPU use with added logic. Better CPU use but complex.
 - **Asynchronous:** Process initiates I/O, keeps running, receives notification on completion. Most efficient, highest complexity.

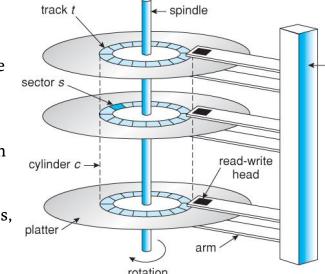
Device Drivers: Two level abstraction separates generic OS interface (top half) from hardware specific operations (bottom half). Issue: Different devices need consistent interface. Solution: Top half handles standard requests (read / write) while bottom half manages specific hardware details and interrupts. I/O Request Lifecycle:

4. Device signals completion via interrupt
5. Bottom half processes interrupt, returns to top half
6. Result returns to application

Persistence: Storage is essential for retaining data without power, unlike RAM. Ideal storage balances capacity, speed and cost. How do we persist data?

Disks: Store data in 512-byte sectors on spinning platters, with read/write heads detecting magnetic domains representing 0s and 1s. Access time is influenced by:

- Seek Time: Time for the read/write heads to reach the correct track.
- Rotational Latency: Delay as the disk rotates to align the target sector under the head.
- Transfer Time: Time to read/write data, depending on RPM and sector density.



Performance metrics:

- Random Access: Involves frequent seeks and rotations, focusing on average seek time and latency.
- Sequential Access: Reads/writes adjacent sectors, minimizing movement and maximizing transfer rate.

SSDs store data by trapping electric charges in flash memory cells, where each cell holds a binary value based on its charge state. Data is organized into pages (for reading) and blocks (for erasing), with a controller managing data placement and wear levelling to extend drive lifespan.

- Pros: Fast read/write speeds, low latency, durability, silent operation, low power, compact.
- Cons: Higher cost, limited write endurance, sudden failure risk, challenging data recovery, smaller max capacity.

I/O Scheduling: Determines the order in which I/O requests are processed to optimize device performance.

- **First-Come, First-Served (FCFS):** Processes requests in the order they arrive. Simple but can cause long waiting times if requests are scattered.
- **Shortest Positioning Time First (SPTF):** Selects the request closest to the disk head, minimizing movement, though it may delay distant requests.
- **SCAN (Elevator):** Moves the disk head in one direction, servicing requests in its path, then reverses. This balances efficiency and fairness by reducing back-and-forth movement.

Performance Metric: Work Conservation ensures the disk remains active when there are pending requests, maximizing utilization.

RAID combines multiple disks to improve storage, performance, and reliability.

- **RAID 0 - Striping:** Splits data across disks, no redundancy. Increases speed but any disk failure results in data loss.
- **RAID 1 - Mirroring:** Copies data to two disks, providing high fault tolerance and faster reads. Halves storage capacity.
- **RAID 4 - Striping with Dedicated Parity:** Data is striped across disks, with one disk storing parity (error-checking data) for recovery if one disk fails. Slower writes due to single parity disk bottleneck.
- **RAID 5 - Distributed Parity:** Data and parity are striped across all disks, allowing recovery from one disk failure and better write performance than RAID 4.

N=#disks, C=capacity, S=sequential throughput, R=random throughput, D=latency of I/O operation

RAID Level	Capacity	Reliability	Read Latency	Write Latency	Seq Read	Seq Write	Rand Read	Rand Write
RAID 0	N * C	Lowest	D	D	N * S	N * S	N * R	N * R
RAID 1	N/2 * C	Highest	D	D	N/2 * S	N/2 * S	N * R	N/2 * R
RAID 4	(N-1) * C	Medium	D	2D	(N-1) * S	(N-1) * S	(N-1) * R	R/2
RAID 5	(N-1) * C	Medium	D	2D	(N-1) * S	(N-1) * S	N * R	N/4 * R

Unix File Systems: Files are abstractions representing persistent byte arrays in storage, enabling data retention beyond process lifetimes.

- **i-nodes:** Data structures holding metadata (size, permissions, block addresses) for a file, identified by an i-node number.
- **Directory Tree (Unix Name Space):** Hierarchical structure where directories map file names to i-node numbers, creating a human-readable path system.
- **File Descriptors:** Non-negative integers assigned by the kernel when a file is opened; used to access files without repeated directory traversals.
- **Hard Links:** Multiple directory entries that reference the same i-node, enabling a single file to have multiple names. Deleting one hard link does not affect others; all hard links must be deleted for the file data to be removed.
- **Symbolic Links:** Files that store a path to another file or directory, allowing links across filesystems and to directories. Symbolic links have different i-node numbers than the target file. If the target file is deleted, symbolic links break, as they rely on the original file path.
- **Permissions:** Control access to files through user/group/world permissions, specified as read (r), write (w), and execute (x).
- **Multi-level Block Indexing:** Uses a mix of direct and indirect pointers to manage large files by addressing blocks hierarchically.

File System Solutions: Locality issues cause slow access when related files are scattered across a disk, leading to frequent seek operations. Solutions focus on keeping related files physically close.

- **Fast File System (FFS):** Uses clusters of consecutive blocks (cylinder groups) to minimize seek times by keeping related structures nearby:

- Structures: Superblock (stores metadata), Allocation Bitmaps (track free/used blocks), Data Blocks (contain file contents), Directories (map filenames to i-node numbers).
- Crash Recovery: FSCK (File System Consistency Check) scans across the entire file system to detect issues such as orphaned inodes, incorrect free block counts, and broken directory links. Issues: time-consuming, requires the file system to be unmounted, cannot always restore files to their original state, often resulting in data loss or orphaned files.

- **Log-Structured File System (LFS):** Writes data sequentially in a log to reduce seek time and optimize for write performance.

- Structures: Segments (log-based storage units for files), Imap (maps i-nodes to storage locations in the log).
- Crash Recovery: Uses checkpointing and roll-forward from the log to restore data after a crash.
- Garbage Collection: Reclaims space by identifying "dead" blocks (stale or deleted data) in segments, compacting live data to free up entire segments for new sequential writes.

- **Journaling:** Maintains a journal to log changes before committing to disk, enabling efficient crash recovery. Faster than FSCK, checkpointing.

- Structures: Journal (logs pending changes to ensure consistency).
- Crash Recovery: Replay journal entries to recover data up to the last committed transaction, minimizing data loss. Checksums: Validates journal entries to detect and prevent corruption.

