

Concurrency: Sequential execution limits resource use and performance. Concurrency interleaves execution of threads or processes to improve utilization and responsiveness.

- **Race Conditions:** Occur when threads access shared resources simultaneously, causing unpredictable results due to interleaved operations.
- **Critical Sections:** Code segments accessing shared resources, requiring protection to prevent race conditions.

Synchronization Primitives: Tools to coordinate concurrent threads, ensuring orderly access to shared resources and preventing race conditions and deadlocks.

- **Locks:** Basic mutual exclusion allowing only one thread in a critical section.
 - Issue: High contention leads to inefficient waiting.
- **Spin Locks:** Threads "spin," repeatedly checking until the lock is free. Use for short critical sections with low contention, as they waste CPU during high contention. Issue: Spinning wastes CPU resources under high contention.
 - Disable Interrupts: Prevents context switches on uniprocessor systems. Fixes context switching but monopolizes the CPU.
 - Load/Store: Loads the lock's current value to check if it's free (0), then stores 1 to claim it. Issue: Not atomic, so race conditions can occur if multiple threads load and store simultaneously.
 - Test-and-Set: Atomically sets the lock's value to 1 and loads the previous value. If the previous value was 0 (unlocked), the thread acquires the lock; otherwise, it retries. Issue: Causes repeated retries (busy-waiting) when locked.
 - Compare-and-Swap (CAS): Atomically loads the lock's current value and compares it with 0. If 0, CAS stores 1 to claim the lock. If not, CAS does nothing, and the thread retries. Issue: Still involves busy-waiting, which wastes CPU.
- **Blocking Locks** (e.g., Mutexes): Threads block (sleep) if the lock is unavailable, reducing CPU waste. Suitable for long critical sections with high contention. Issue: Can cause deadlocks if threads hold resources while waiting for others; lacks conditional waiting.
- **Condition Variables:** Used alongside blocking locks (like mutexes) to allow threads to wait for specific conditions and signal others when conditions change.
 - `wait(mutex)`: Releases the mutex and waits, reacquiring it on wakeup.
 - `signal() / broadcast()`: Wakes up one or all waiting threads.
 - Issue: Spurious wakeups require rechecking the condition; doesn't manage access to multiple resources.
 - Bounded Buffer with Condition Variable: Uses two condition variables, `notEmpty` and `notFull`, for producer and consumer threads. Producers wait on `notFull` when the buffer is full; consumers wait on `notEmpty` when it's empty.
 - Producer/Consumer with Two Condition Variables: `notEmpty` signals consumers when items are available; `notFull` signals producers when there's space.
- **Semaphores:** Control access to multiple instances of a resource. A non-negative integer count represents available resources.
 - `acquire()`: Decrements count if positive; blocks if zero.
 - `release()`: Increments count; may wake waiting threads.
 - Bounded Buffer with Semaphores: Use a semaphore empty for free slots and full for filled slots, alongside a mutex to control access. Producers `wait(empty)` before adding; consumers `wait(full)` before removing.
- **Equivalence:** Synchronization primitives are functionally equivalent by adapting usage. A binary semaphore acts as a mutex by setting its count to 0 (`acquire`) and 1 (`release`) for mutual exclusion. Conversely, a mutex and condition variable mimic a counting semaphore by using a resource count and signalling waiting threads, enabling controlled access.

Deadlock: Occurs when threads are blocked, each waiting for resources held by others. 4 Conditions must be met to reach a deadlock, if any one condition is prevented, the deadlock is broken:

1. Mutual Exclusion: Only one thread can use a resource at a time.
2. Hold and Wait: Threads hold resources while waiting for others.
3. No Pre-emption: Resources cannot be forcibly taken away.
4. Circular Wait: Circular chain of threads each waiting for resources held by the next.

Avoidance:

- Wait-Free Algorithms: Ensure every operation completes in a finite number of steps, removing the need to wait.
- Atomic Lock Acquisition: Require threads to acquire all needed resources atomically.
- Pre-emption: Allow resources to be forcibly taken away.
- Resource Ordering: Impose an order on resource acquisition to prevent circular wait.
- Banker's Algorithm:** Checks if granting a resource can be done safely without causing a deadlock. Parameters include:
 - Safe State: A state where the system can allocate resources to each process in some order and still avoid deadlock.
 - Available Resources: The number of available resources of each type.
 - Maximum Demand: The maximum number of resources a process may need.
 - Allocated Resources: The number of resources currently allocated to a process.
 - Need: The remaining resources that a process may still request.
- a. Check Need: Ensure the request doesn't exceed what the process still needs.
- b. Check Availability: Confirm enough resources are available to fulfill the request.
- c. Simulate Allocation: Temporarily grant the resources and verify if the system stays safe.
- d. Decision: If safe, grant the request. If not, deny and make the process wait.

Given: 3 people (Alice, Bob, Charlie), 3 fruits (Apples, Bananas, Oranges), Available: 2 1 0
Alice requests 1 1 0 (Need: 2 2 0, Allocated: 1 0 1) → Granted, no one starves
Charlie requests 1 0 1 (Need: 1 1 1, Allocated: 0 1 0) → Denied, exceeds Need

I/O Devices: Users need to be able interact with OS via external devices. OS Needs standardised way to interact with diverse hardware having different speeds and protocols.

- **Canonical devices** solve hardware diversity through standardized registers (status: reports device state, command: sends instructions, data: transfers information). Driver translates between standard interface and hardware. Device interaction happens through:
 - **Polling:** CPU repeatedly checks device status. Issue: Wastes CPU cycles - good for fast devices or predictable timing.
 - **Interrupts:** Device signals CPU when ready. Issue: Polling inefficient. Solution: CPU free to do other work while waiting.
 - **Direct Memory Access (DMA):** Issue: CPU wastes cycles copying data, so DMA transfers data between memory and devices independently. CPU only sets up transfer then continues work; gets interrupted when done.
- **Patterns:** What does the process do while waiting on I/O?
 - **Blocking:** Process waits until I/O completes, simple but idle CPU.
 - **Non Blocking:** Process starts I/O, keeps running, checks status periodically. Improves CPU use with added logic. Better CPU use but complex.
 - **Asynchronous:** Process initiates I/O, keeps running, receives notification on completion. Most efficient, highest complexity.


Device Drivers: Two level abstraction separates generic OS interface (top half) from hardware specific operations (bottom half). Issue: Different devices need consistent interface. Solution: Top half handles standard requests (read / write) while bottom half manages specific hardware details and interrupts. **I/O Request Lifecycle:**

1. Application sends request through top half
2. Driver converts to device commands in bottom half
3. Device executes operation
4. Device signals completion via interrupt
5. Bottom half processes interrupt, returns to top half
6. Result returns to application

Persistence: Storage is essential for retaining data without power, unlike RAM. Ideal storage balances capacity, speed and cost. How do we persist data?

Disk: Store data in 512-byte sectors on spinning platters, with read/write heads detecting magnetic domains representing 0s and 1s. Access time is influenced by:

- Seek Time: Time for the read/write heads to reach the correct track.
- Rotational Latency: Delay as the disk rotates to align the target sector under the head.
- Transfer Time: Time to read/write data, depending on RPM and sector density.



SSD: Store data by trapping electric charges in flash memory cells, where each cell holds a binary value based on its charge state. Data is organized into pages (for reading) and blocks (for erasing), with a controller managing data placement and wear levelling to extend drive lifespans.

- Pros: Fast read/write speeds, low latency, durability, silent operation, low power, compact.
- Cons: Higher cost, limited write endurance, sudden failure risk, challenging data recovery, smaller max capacity.

I/O Scheduling: Determines the order in which I/O requests are processed to optimize device performance.

- **First-Come, First-Served (FCFS):** Processes requests in the order they arrive. Simple but can cause long waiting times if requests are scattered.
- **Shortest Positioning Time First (SPTF):** Selects the request closest to the disk head, minimizing movement, though it may delay distant requests.
- **SCAN (Elevator):** Moves the disk head in one direction, servicing requests in its path, then reverses. This balances efficiency and fairness by reducing back-and-forth movement.

Performance Metric: Work Conservation ensures the disk remains active when there are pending requests, maximizing utilization.

RAID: Combines multiple disks to improve storage, performance, and reliability.

- **RAID 0 - Striping:** Splits data across disks, no redundancy. Increases speed but any disk failure results in data loss.
- **RAID 1 - Mirroring:** Copies data to two disks, providing high fault tolerance and faster reads. Halves storage capacity.
- **RAID 4 - Striping with Dedicated Parity:** Data is striped across disks, with one disk storing parity (error-checking data) for recovery if one disk fails. Slower writes due to single parity disk bottleneck.
- **RAID 5 - Distributed Parity:** Data and parity are striped across all disks, allowing recovery from one disk failure and better write performance than RAID 4.

N=disks, C=capacity, S=sequential throughput, R=random throughput, D=latency of I/O operation

| RAID Level | Capacity | Reliability | Read Latency | Write Latency | Seq Read | Seq Write | Rand Read | Rand Write |
|------------|-----------|-------------|--------------|---------------|-----------|-----------|-----------|------------|
| RAID 0 | N * C | Lowest | D | D | N * S | N * S | N * R | N * R |
| RAID 1 | N/2 * C | Highest | D | D | N/2 * S | N/2 * S | N * R | N/2 * R |
| RAID 4 | (N-1) * C | Medium | D | 2D | (N-1) * S | (N-1) * S | (N-1) * R | R/2 |
| RAID 5 | (N-1) * C | Medium | D | 2D | (N-1) * S | (N-1) * S | N * R | N/4 * R |

Unix File Systems: Files are abstractions representing persistent byte arrays in storage, enabling data retention beyond process lifetimes.

- **i-nodes:** Data structures holding metadata (size, permissions, block addresses) for a file, identified by an i-node number.
- **Directory Tree (Unix Name Space):** Hierarchical structure where directories map file names to i-node numbers, creating a human-readable path system.
- **File Descriptors:** Non-negative integers assigned by the kernel when a file is opened; used to access files without repeated directory traversals.
- **Hard Links:** Multiple directory entries that reference the same i-node, enabling a single file to have multiple names. Deleting one hard link does not affect others; all hard links must be deleted for the file data to be removed.
- **Symbolic Links:** Files that store a path to another file or directory, allowing links across filesystems and to directories. Symbolic links have different i-node numbers than the target file. If the target file is deleted, symbolic links break, as they rely on the original file path.
- **Permissions:** Control access to files through user/group/world permissions, specified as read (r), write (w), and execute (x).
- **Multi-level Block Indexing:** Uses a mix of direct and indirect pointers to manage large files by addressing blocks hierarchically.

File System Solutions: Locality issues cause slow access when related files are scattered across a disk, leading to frequent seek operations. Solutions focus on keeping related files physically close.

- **Fast File System (FFS):** Uses clusters of consecutive blocks (cylinder groups) to minimize seek times by keeping related structures nearby:
 - Structures: Superblock (stores metadata), Allocation Bitmaps (track free/used blocks), Data Blocks (contain file contents), Directories (map filenames to i-node numbers).
 - Crash Recovery: FSCK (File System Consistency Check) scans across the entire file system to detect issues such as orphaned inodes, incorrect free block counts, and broken directory links. Issues: time-consuming, requires the file system to be unmounted, cannot always restore files to their original state, often resulting in data loss or orphaned files.
- **Log-Structured File System (LFS):** Writes data sequentially in a log to reduce seek time and optimize for write performance.
 - Structures: Segments (log-based storage units for files), Imap (maps i-nodes to storage locations in the log).
 - Crash Recovery: Uses checkpointing and roll-forward from the log to restore data after a crash.
 - Garbage Collection: Reclaims space by identifying "dead" blocks (stale or deleted data) in segments, compacting live data to free up entire segments for new sequential writes.
- **Journaling:** Maintains a journal to log changes before committing to disk, enabling efficient crash recovery. Faster than FSCK, checkpointing.
 - Structures: Journal (logs pending changes to ensure consistency).
 - Crash Recovery: Replay journal entries to recover data up to the last committed transaction, minimizing data loss. Checksums: Validates journal entries to detect and prevent corruption.

