

human ends having to do with love and friendship. They need the ability to form family ties, friendships, and special relations with neighbours. They also need codes of conduct. And how could they have all these things without virtues such as loyalty, fairness, kindness, and in certain circumstances obedience? . . .

Thus the structure of the derivation is the same whether we derive an evaluation of the roots of a particular tree or the action of a particular human being. The meaning of the words 'good' and 'bad' is not different when used of features of plants on the one hand and humans on the other, but is rather the same as applied, in judgements of natural goodness and defect, in the case of all living things.

### Philippa Foot: Natural Goodness

1. What is the difference, according to Foot, between "natural goodness" and "secondary goodness"?
2. Consider the sentence "Cats are four-legged." Foot points out that such a sentence doesn't claim that *all* cats have four legs. Yet it also claims more than the sentence "Some cats are four-legged." What, then, does such a sentence mean?
3. Foot says that a creature is *defective* if it does not possess the traits that "fit it to survive" in its "natural habitat." It follows that a deer is defective if it is unable to run very fast, even if that deer is in a zoo and doesn't need to run from any predators. Is this a plausible claim?
4. What is Foot's objection to the view that a good life is just a happy life? Do you find this objection persuasive?
5. Foot says that the words "good" and "bad" mean the same thing when applied to plants as when applied to humans. What does she mean by this claim? Do you agree with her?

## Extreme and Restricted Utilitarianism

J. J. C. Smart

J. J. C. Smart here introduces us to two forms of utilitarianism, extreme and restricted, which nowadays go by the name of act and rule utilitarianism, respectively. Extreme utilitarianism makes the morality of actions depend entirely on their results, requiring us always to do that action, among all available to us at the time, which will yield the greatest overall happiness. Restricted (i.e., rule) utilitarianism, by contrast, assesses the rightness of actions based on whether they adhere to rules that, if embraced by all, would yield optimal happiness.

Smart himself endorses extreme utilitarianism, and presents a criticism of restricted utilitarianism that has been very influential. Suppose that we are faced with a situation in which we must break a moral rule in order to maximize happiness. Since the goal of such rules, from a utilitarian standpoint, is to maximize the world's happiness, the only rational thing to do is to break that rule. And yet restricted utilitarians insist that we honor the rule, even when doing so undermines the ultimate utilitarian goal. That is irrational, says Smart, and gives us decisive reason to prefer the extreme version of utilitarianism.

## 1

Utilitarianism is the doctrine that the rightness of actions is to be judged by their consequences. What do we mean by "actions" here? Do we mean particular actions or do we mean classes of actions? According to which way we interpret the word "actions" we get two different theories, both of which merit the appellation "utilitarian."

(1) If by "actions" we mean particular individual actions we get the sort of doctrine held by Bentham, Sidgwick, and Moore. According to this doctrine we test individual actions by their consequences, and general rules, like "keep promises," are mere rules of thumb which we use only to avoid the necessity of estimating the probable consequences of our actions at every step. The rightness or wrongness of keeping a promise on a particular occasion depends only on the goodness or badness of the consequences of keeping or of breaking the promise on that particular occasion. Of course part of the consequences of breaking the promise, and a part to which we will normally ascribe decisive importance, will be the weakening of faith in the institution of promising. However, if the goodness of the consequences of breaking the rule is *in toto* greater than the goodness of the consequences of keeping it, then we must break the rule, irrespective of whether the goodness of the consequences of *everybody's* obeying the rule is or is not greater than the consequences of *everybody's* breaking it. To put it shortly, rules do not matter, save *per accidens* as rules of thumb and as *de facto* social institutions with which the utilitarian has to reckon when estimating consequences. I shall call this doctrine "extreme utilitarianism."

(2) A more modest form of utilitarianism has recently become fashionable. The doctrine is to be found in Toulmin's book *The Place of Reason in Ethics*, in Nowell-Smith's *Ethics* (though I think Nowell-Smith has qualms), in John Austin's *Lectures on Jurisprudence* (Lecture II), and even in J. S. Mill, if Urmson's interpretation of him is correct (*Philosophical Quarterly*, vol. 3, pp. 33–9, 1953). Part of its charm is that it appears to resolve the dispute in moral philosophy between intuitionists and utilitarians in a way which is very neat. The above philosophers hold, or seem to hold, that moral rules are more than rules of thumb. In general the rightness of an action is *not* to be tested by evaluating its consequences but only by considering whether or not it falls under a certain rule. Whether the rule is to be considered an acceptable moral rule, is, however, to be decided by considering the consequences of adopting the rule.

Broadly, then, actions are to be tested by rules and rules by consequences. The only cases in which we must test an individual action directly by its consequences are (a) when the action comes under two different rules, one of which enjoins it and one of which forbids it, and (b) when there is no rule whatever that governs the given case. I shall call this doctrine "restricted utilitarianism."

## 2

For an extreme utilitarian moral rules are rules of thumb. In practice the extreme utilitarian will mostly guide his conduct by appealing to the rules ("do not lie," "do not break promises," etc.) of common sense morality. This is not because there is anything sacrosanct in the rules themselves but because he can argue that probably he will most often act in an extreme utilitarian way if he does not think as a utilitarian. For one thing, actions have frequently to be done in a hurry. Imagine a man seeing a person drowning. He jumps in and rescues him. There is no time to reason the matter out, but usually this will be the course of action which an extreme utilitarian would recommend if he did reason the matter out. If, however, the man drowning had been drowning in a river near Berchtesgaden in 1938, and if he had had the well known black forelock and moustache of Adolf Hitler, an extreme utilitarian would, if he had time, work out the probability of the man's being the villainous dictator, and if the probability were high enough he would, on extreme utilitarian grounds, leave him to drown. The rescuer, however, has not time. He trusts to his instincts and dives in and rescues the man. And this trusting to instincts and to moral rules can be justified on extreme utilitarian grounds. Furthermore, an extreme utilitarian who knew that the drowning man was Hitler would nevertheless praise the rescuer, not condemn him. For by praising the man he is strengthening a courageous and benevolent disposition of mind, and in general this disposition has great positive utility. (Next time, perhaps, it will be Winston Churchill that the man saves!) We must never forget that an extreme utilitarian may praise actions which he knows to be wrong. Saving Hitler was wrong, but it was a member of a class of actions which are generally right, and the motive to do actions of this class is in general an optimistic one. In considering questions of praise and blame it is not the expediency of the praised or blamed action that is at issue, but the expediency of the praise. It can be expedient to praise an inexpedient action and inexpedient to praise an expedient one.

Lack of time is not the only reason why an extreme utilitarian may, on extreme utilitarian principles, trust to rules of common sense morality. He knows that in particular cases where his own interests are involved his calculations are likely to be biased in his own favor. Suppose that he is unhappily married and is deciding whether to get divorced. He will in all probability greatly exaggerate his own unhappiness (and possibly his wife's) and greatly underestimate the harm done to his children by the break up of the family. He will probably also underestimate the likely harm done by the weakening of the general faith in marriage vows. So probably he will come to the correct extreme utilitarian conclusion if he does not in this instance think as an extreme utilitarian but trusts to common sense morality.

There are many more and subtle points that could be made in connection with the relation between extreme utilitarianism and the morality of common sense. All those that I have just made and many more will be found in Book IV Chapters 3–5 of Sidgwick's *Methods of Ethics*. I think that this book is the best book ever written on ethics, and that these chapters are the best chapters of the book. As they occur so near the end of a very long book they are unduly neglected. I refer the reader, then, to Sidgwick for the classical exposition of the relation between (extreme) utilitarianism and the morality of common sense. One further point raised by Sidgwick in this connection is whether an (extreme) utilitarian ought on (extreme) utilitarian principles to propagate (extreme) utilitarianism among the public. As most people are not very philosophical and not good at empirical calculations, it is probable that they will most often act in an extreme utilitarian way if they do not try to think as extreme utilitarians. We have seen how easy it would be to misapply the extreme utilitarian criterion in the case of divorce. Sidgwick seems to think it quite probable that an extreme utilitarian should not propagate his doctrine too widely. However, the great danger to humanity comes nowadays on the plane of public morality—not private morality. There is a greater danger to humanity from the hydrogen bomb than from an increase of the divorce rate, regrettable though that might be, and there seems no doubt that extreme utilitarianism makes for good sense in international relations. When France walked out of the United Nations because she did not wish Morocco discussed, she said that she was within her rights because Morocco and Algiers are part of her metropolitan territory and nothing to do with UN. This was clearly a legalistic if not superstitious argument. We should not be concerned with the so-called “rights” of

France or any other country but with whether the cause of humanity would best be served by discussing Morocco in UN. (I am not saying that the answer to this is “Yes.” There are good grounds for supposing that more harm than good would come by such a discussion.) I myself have no hesitation in saying that on extreme utilitarian principles we ought to propagate extreme utilitarianism as widely as possible. But Sidgwick had respectable reasons for suspecting the opposite.

The extreme utilitarian, then, regards moral rules as rules of thumb and as sociological facts that have to be taken into account when deciding what to do, just as facts of any other sort have to be taken into account. But in themselves they do not justify any action.

### 3

The restricted utilitarian regards moral rules as more than rules of thumb for short-circuiting calculations of consequences. Generally, he argues, consequences are not relevant at all when we are deciding what to do in a particular case. In general, they are relevant only to deciding what rules are good reasons for acting in a certain way in particular cases. This doctrine is possibly a good account of how the modern unreflective twentieth century Englishman often thinks about morality, but surely it is monstrous as an account of how it is most rational to think about morality. Suppose that there is a rule *R* and that in 99% of cases the best possible results are obtained by acting in accordance with *R*. Then clearly *R* is a useful rule of thumb; if we have not time or are not impartial enough to assess the consequences of an action it is an extremely good bet that the thing to do is to act in accordance with *R*. But is it not monstrous to suppose that if we *have* worked out the consequences and if we have perfect faith in the impartiality of our calculations, and if we *know* that in this instance to break *R* will have better results than to keep it, we should nevertheless obey the rule? Is it not to erect *R* into a sort of idol if we keep it when breaking it will prevent, say, some avoidable misery? Is not this a form of superstitious rule-worship (easily explicable psychologically) and not the rational thought of a philosopher?

The point may be made more clearly if we consider Mill's comparison of moral rules to the tables in the nautical almanac (*Utilitarianism*, Everyman Edition, pp. 22–3). This comparison of Mill's is adduced by Urmson as evidence that Mill was a restricted utilitarian, but I do not think that it will bear this interpretation at all. (Though I quite agree with Urmson that

many other things said by Mill are in harmony with restricted rather than extreme utilitarianism. Probably Mill had never thought very much about the distinction and was arguing for utilitarianism, restricted or extreme, against other and quite non-utilitarian forms of moral argument.) Mill says: "Nobody argues that the art of navigation is not founded on astronomy, because sailors cannot wait to calculate the Nautical Almanac. Being rational creatures, they go out upon the sea of life with their minds made up on the common questions of right and wrong, as well as on many of the far more difficult questions of wise and foolish. . . . Whatever we adopt as the fundamental principle of morality, we require subordinate principles to apply it by." Notice that this is, as it stands, only an argument for subordinate principles as rules of thumb. The example of the nautical almanac is misleading because the information given in the almanac is in all cases the same as the information one would get if one made a long and laborious calculation from the original astronomical data on which the almanac is founded. Suppose, however, that astronomy were different. Suppose that the behavior of the sun, moon and planets was very nearly as it is now, but that on rare occasions there were peculiar irregularities and discontinuities, so that the almanac gave us rules of the form "in 99% of cases where the observations are such and such you can deduce that your position is so and so." Furthermore, let us suppose that there were methods which enabled us, by direct and laborious calculation from the original astronomical data, not using the rough and ready tables of the almanac, to get our correct position in 100% of cases. Seafarers might use the almanac because they never had time for the long calculations and they were content with a 99% chance of success in calculating their positions. Would it not be absurd, however, if they *did* make the direct calculation, and finding that it disagreed with the almanac calculation, nevertheless they ignored it and stuck to the almanac conclusion? Of course the case would be altered if there were a high enough probability of making slips in the direct calculation: then we might stick to the almanac result, liable to error though we knew it to be, simply because the direct calculation would be open to error for a different reason, the fallibility of the computer. This would be analogous to the case of the extreme utilitarian who abides by the conventional rule against the dictates of his utilitarian calculations simply because he thinks that his calculations are probably affected by personal bias. But if the navigator were sure of his direct calculations would he not be foolish to abide by his almanac? I conclude, then, that if we change our suppositions about astronomy and the almanac (to which

there are no exceptions) to bring the case into line with that of morality (to whose rules there are exceptions), Mill's example loses its appearance of supporting the restricted form of utilitarianism. Let me say once more that I am not here concerned with how ordinary men think about morality but with how they ought to think. We could quite well imagine a race of sailors who acquired a superstitious reverence for their almanac, even though it was only right in 99% of cases, and who indignantly threw overboard any man who mentioned the possibility of a direct calculation. But would this behavior of the sailors be rational?

Let us consider a much discussed sort of case in which the extreme utilitarian might go against the conventional moral rule. I have promised to a friend, dying on a desert island from which I am subsequently rescued, that I will see that his fortune (over which I have control) is given to a jockey club. However, when I am rescued I decide that it would be better to give the money to a hospital, which can do more good with it. It may be argued that I am wrong to give the money to the hospital. But why? (a) The hospital can do more good with the money than the jockey club can. (b) The present case is unlike most cases of promising in that no one except me knows about the promise. In breaking the promise I am doing so with complete secrecy and am doing nothing to weaken the general faith in promises. That is, a factor, which would normally keep the extreme utilitarian from promise breaking even in otherwise unoptimal cases, does not at present operate. (c) There is no doubt a slight weakening in my own character as an habitual promise keeper, and moreover psychological tensions will be set up in me every time I am asked what the man made me promise him to do. For clearly I shall have to say that he made me promise to give the money to the hospital, and, since I am an habitual truth teller, this will go very much against the grain with me. Indeed I am pretty sure that in practice I myself would keep the promise. But we are not discussing what my moral habits would probably make me do; we are discussing what I ought to do. Moreover, we must not forget that even if it would be most rational of me to give the money to the hospital it would also be most rational of you to punish or condemn me if you did, most improbably, find out the truth (e.g. by finding a note washed ashore in a bottle). Furthermore, I would agree that though it was most rational of me to give the money to the hospital it would be most rational of you to condemn me for it. We revert again to Sidgwick's distinction between the utility of the action and the utility of the praise of it.



Many such issues are discussed by A. K. Stout<sup>1</sup> . . . It will be useful . . . to consider one . . . example that he gives. Suppose that during hot weather there is an edict that no water must be used for watering gardens. I have a garden and I reason that most people are sure to obey the edict, and that as the amount of water that I use will be by itself negligible no harm will be done if I use the water secretly. So I do use the water, thus producing some lovely flowers which give happiness to various people. Still, you may say, though the action was perhaps optimific, it was unfair and wrong.

There are several matters to consider. Certainly my action should be condemned. We revert once more to Sidgwick's distinction. A right action may be rationally condemned. Furthermore, this sort of offense is normally found out. If I have a wonderful garden when everybody else's is dry and brown there is only one explanation. So if I water my garden I am weakening my respect for law and order, and as this leads to bad results an extreme utilitarian would agree that I was wrong to water the garden. Suppose now that the case is altered and that I can keep the thing secret: there is a secluded part of the garden where I grow flowers which I give away anonymously to a home for old ladies. Are you still so sure that I did the wrong thing by watering my garden? However, this is still a weaker case than that of the hospital and the jockey club. There will be tensions set up within myself: my secret knowledge that I have broken the rule will make it hard for me to exhort others to keep the rule. These psychological ill effects in myself may be not inconsiderable: directly and indirectly they may lead to harm which is at least of the same order as the happiness that the old ladies get from the flowers. You can see that on an extreme utilitarian view there are two sides to the question.

So far I have been considering the duty of an extreme utilitarian in a predominantly non-utilitarian society. . . . I now pass on to a type of case which may be thought to be the trump card of restricted utilitarianism. Consider the rule of the road. It may be said that since all that matters is that everyone should do the same it is indifferent which rule we have, "go on the left hand side" or "go on the right hand side." Hence the only *reason* for going on the left hand side in British countries is that this is the rule. Here the rule does seem to be a reason, in itself, for acting in a certain way. I wish to argue against this. The rule in itself is not a reason for our actions. We would be perfectly justified in going on the right hand side if (a) we knew that the rule was to go on the left hand side, and (b) we were in a

country peopled by super-anarchists who always on principle did the opposite of what they were told. This shows that the rule does not give us a reason for acting so much as an indication of the probable actions of others, which helps us to find out what would be our own most rational course of action. If we are in a country not peopled by anarchists, but by non-anarchist extreme Utilitarians, we expect, other things being equal, that they will keep rules laid down for them. Knowledge of the rule enables us to predict their behavior and to harmonize our own actions with theirs. The rule "keep to the left hand side," then, is not a logical *reason* for action but an anthropological *datum* for planning actions.

I conclude that in every case if there is a rule *R* the keeping of which is in general optimific, but such that in a special sort of circumstances the optimific behavior is to break *R*, then in these circumstances we should break *R*. Of course we must consider all the less obvious effects of breaking *R*, such as reducing people's faith in the moral order, before coming to the conclusion that to break *R* is right: in fact we shall rarely come to such a conclusion. Moral rules, on the extreme utilitarian view, are rules of thumb only, but they are not bad rules of thumb. But if we *do* come to the conclusion that we should break the rule and if we have weighed in the balance our own fallibility and liability to personal bias, what good reason remains for keeping the rule? I can understand "it is optimific" as a reason for action, but why should "it is a member of a class of actions which are usually optimific" or "it is a member of a class of actions which as a class are more optimific than any alternative general class" be a good reason? You might as well say that a person ought to be picked to play for Australia just because all his brothers have been, or that the Australian team should be composed entirely of the Harvey family because this would be better than composing it entirely of any other family. The extreme utilitarian does not appeal to artificial feelings, but only to our feelings of benevolence, and what better feelings can there be to appeal to? Admittedly we can have a pro-attitude to anything, even to rules, but such artificially begotten pro-attitudes smack of superstition. Let us get down to realities, human happiness and misery, and make these the objects of our pro-attitudes and anti-attitudes.

The restricted utilitarian might say that he is talking only of *morality*, not of such things as rules of the road. I am not sure how far this objection, if valid, would affect my argument, but in any case I would reply that as a philosopher I conceive of ethics as the study of how it would be *most rational* to act. If my opponent wishes to restrict the word "morality" to a

1 "But Suppose Everybody Did the Same?," *Australasian Journal of Philosophy* vol. 32, pp. 1-29.

narrower use he can have the word. The fundamental question is the question of rationality of action *in general*. Similarly if the restricted utilitarian were to appeal to ordinary usage and say "it might be most rational to leave Hitler to drown but it would surely not be *wrong* to rescue him," I should again let him have the words "right" and "wrong" and should stick to "rational" and "irrational." We already saw that it would be rational to praise Hitler's rescuer, even though it would have been most rational not to have rescued Hitler. In ordinary language, no doubt, "right" and "wrong" have not only the meaning "most rational to do" and "not most rational to do" but also have the meaning "praiseworthy" and "not praiseworthy." Usually to the utility of an action corresponds utility of praise of it, but as we saw, this is not always so. Moral language could thus do with tidying up, for example by reserving "right" for "most rational" and "good" as an epithet of praise for the motive from which the action sprang. It would be more becoming in a philosopher to try to iron out illogicalities in moral language and to make suggestions for its reform than to use it as a court of appeal whereby to perpetuate confusions.

### J.J.C. Smart: Extreme and Restricted Utilitarianism

1. Extreme utilitarians hold that moral rules are "mere rules of thumb," while restricted utilitarians think that the rightness of an action is determined by "whether or not it falls under a certain rule." What exactly is the difference between these two views of moral rules?
2. According to Smart, we should sometimes praise actions we know to be wrong, and condemn actions we know to be right. Why does he think this? Do you agree?
3. Smart claims that restricted utilitarianism amounts to "superstitious rule-worship." What does he mean by this? Do you think he is correct?
4. Extreme utilitarians think that we ought to break our promises whenever doing so will bring about a better outcome than keeping the promise. Are we obligated to keep our promises, even if we could bring about more good by breaking them?
5. In his example of the gardener, Smart suggests that it is sometimes morally acceptable to secretly break the law, even if the law is perfectly just. Do you agree with him?

## The Good Will and the Categorical Imperative

*Immanuel Kant*

Immanuel Kant (1724–1804) was the greatest German philosopher who ever lived. In this excerpt from his *Groundwork of the Metaphysics of Morals*, Kant introduces two key elements of his moral philosophy. According to Kant, the first of these, the *good will*, is the only thing possessed of unconditional value: it is valuable in its own right, in every possible circumstance. The good will is the steady commitment to do our duty for its own sake. Our actions possess moral worth if, but only if, they are prompted by the good will.

The second important element is the *categorical imperative*, Kant's term for a requirement of reason that applies to us regardless of what we care about. Moral requirements are categorical imperatives—we must, for instance, sometimes give help to others in need, even if we don't want to, and even if such help gets us nothing that we care about. Kant believed that moral action is rational action. Each of us has a compelling reason to obey morality, even when doing so only frustrates our deepest desires.

Kant here sets out two tests for morally acceptable action. The first says that actions are morally acceptable only when the principles that inspire them can be acted on by everyone consistently. The second requires us to treat humanity always as an end in itself, and

From *Groundwork of the Metaphysics of Morals*, trans. Mary Gregor (1998). Reprinted with the permission of Cambridge University Press.